

Explanation

1. Load the MNIST dataset and pre-process the data:

- Load the MNIST dataset using a library such as TensorFlow or PyTorch
- Normalize the pixel values to be in the range [0, 1]
- Split the data into training and testing sets

2. Define the MLP model:

- Initialize an empty sequential model using a library such as TensorFlow's keras or PyTorch's nn module
- Add an input layer to flatten the 28x28 images into 1D arrays
- Add one or more dense hidden layers with a specified number of units and activation function (e.g. ReLU)
- Add an output layer with 10 units and a softmax activation function

3. Compile the model:

- Specify the optimizer (e.g. Adam or SGD) to use for training
- Specify the loss function (e.g. cross-entropy loss for multi-class classification)
- Specify any evaluation metrics to track during training (e.g. accuracy)

4. Train the model:

- Train the model on the training data for a specified number of epochs
- Optionally, use validation data to evaluate the model during training and early stop if the performance on the validation data doesn't improve

5. Evaluate the model:

- Evaluate the model on the test data and report the final accuracy or other evaluation metrics

6. Make predictions:

- Use the trained model to make predictions on new data and obtain the predicted values
- Convert the predicted probabilities into predicted class labels by selecting the class with the highest predicted probability